

Salman_NUS_Graded Assignment 7.1: Improved Digit Generation with VAEs

```

import tensorflow as tf
from tensorflow.keras.layers import Input, Dense, Flatten, Reshape, Layer
from tensorflow.keras.models import Model
import numpy as np
import matplotlib.pyplot as plt

# Load and preprocess MNIST data
(x_train, _), (_, _) = tf.keras.datasets.mnist.load_data()
x_train = x_train.astype('float32') / 255.
x_train = x_train.reshape((len(x_train), 28, 28, 1))

# Network parameters
input_shape = (28, 28, 1)
batch_size = 128
latent_dim = 2
intermediate_dim = 512
epochs = 5

class Sampling(Layer):
    def call(self, inputs):
        z_mean, z_log_var = inputs
        batch = tf.shape(z_mean)[0]
        dim = tf.shape(z_mean)[1]
        epsilon = tf.random.normal(shape=(batch, dim))
        return z_mean + tf.exp(0.5 * z_log_var) * epsilon

# Encoder
encoder_inputs = Input(shape=input_shape)
x = Flatten()(encoder_inputs)
x = Dense(intermediate_dim, activation='relu')(x)
z_mean = Dense(latent_dim, name='z_mean')(x)
z_log_var = Dense(latent_dim, name='z_log_var')(x)
z = Sampling()([z_mean, z_log_var])
encoder = Model(encoder_inputs, [z_mean, z_log_var, z], name='encoder')

# Decoder
decoder_inputs = Input(shape=(latent_dim,))
x = Dense(intermediate_dim, activation='relu')(decoder_inputs)
x = Dense(784, activation='sigmoid')(x)
decoder_outputs = Reshape((28, 28, 1))(x)
decoder = Model(decoder_inputs, decoder_outputs, name='decoder')

class VAE(tf.keras.Model):
    def __init__(self, encoder, decoder, **kwargs):
        super().__init__(**kwargs)
        self.encoder = encoder
        self.decoder = decoder
        self.total_loss_tracker = tf.keras.metrics.Mean(name='total_loss')
        self.reconstruction_loss_tracker = tf.keras.metrics.Mean(name='reconstruction_loss')
        self.kl_loss_tracker = tf.keras.metrics.Mean(name='kl_loss')

    @property
    def metrics(self):
        return [
            self.total_loss_tracker,
            self.reconstruction_loss_tracker,
            self.kl_loss_tracker
        ]

    def train_step(self, data):
        with tf.GradientTape() as tape:
            # Encoder output
            z_mean, z_log_var, z = self.encoder(data)
            # Reconstruction
            reconstruction = self.decoder(z)

            # Flatten input and reconstruction for binary crossentropy
            flat_data = tf.reshape(data, (-1, 784))
            flat_reconstruction = tf.reshape(reconstruction, (-1, 784))

            # Reconstruction loss (fixed axis handling)
            reconstruction_loss = tf.reduce_mean(
                tf.keras.losses.binary_crossentropy(
                    flat_data,
                    flat_reconstruction
                )
            ) * 784

            # KL divergence loss
            kl_loss = -0.5 * tf.reduce_mean(
                tf.reduce_sum(
                    1 + z_log_var - tf.square(z_mean) - tf.exp(z_log_var),
                    axis=1
                )
            )

            # Total loss
            total_loss = reconstruction_loss + kl_loss

            # Compute gradients
            grads = tape.gradient(total_loss, self.trainable_weights)
            self.optimizer.apply_gradients(zip(grads, self.trainable_weights))

            # Update metrics
            self.total_loss_tracker.update_state(total_loss)
            self.reconstruction_loss_tracker.update_state(reconstruction_loss)
            self.kl_loss_tracker.update_state(kl_loss)

        return {
            "loss": self.total_loss_tracker.result(),
            "reconstruction_loss": self.reconstruction_loss_tracker.result(),
            "kl_loss": self.kl_loss_tracker.result()
        }

    def call(self, inputs):
        z_mean, z_log_var, z = self.encoder(inputs)
        reconstruction = self.decoder(z)
        return reconstruction

```

```

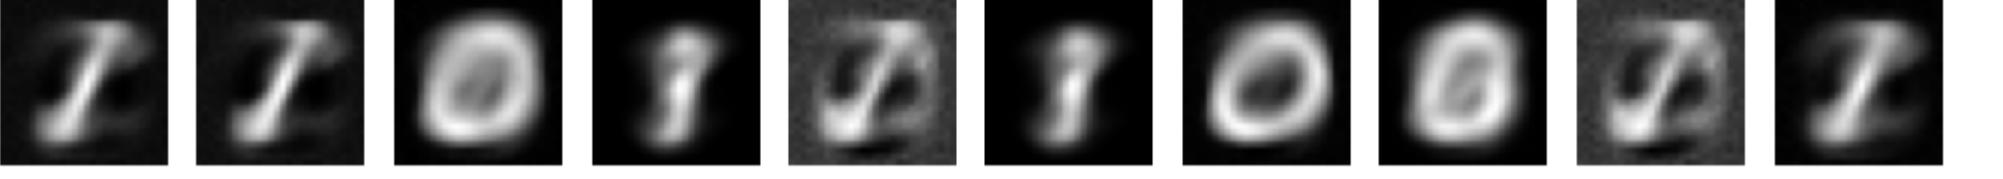
7/5/25, 1:52 PM
# Create and compile VAE
print("Creating VAE...")
vae = VAE(encoder, decoder)
vae.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=1e-4))

# Train VAE
print("Training VAE...")
vae.fit(x_train,
        epochs=epochs,
        batch_size=batch_size,
        shuffle=True)

def generate_digits(decoder, n=10):
    """Generate n new digits using the decoder"""
    z_sample = np.random.normal(size=(n, latent_dim))
    x_decoded = decoder.predict(z_sample)
    return x_decoded

# After training the VAE, you can use it like this:
print("\nGenerating new digits...")
new_digits = generate_digits(decoder, 10) # Generates 10 new digits

# To visualize the generated digits
plt.figure(figsize=(10, 2))
for i in range(10):
    plt.subplot(1, 10, i + 1)
    plt.imshow(new_digits[i].reshape(28, 28), cmap='gray')
    plt.axis('off')
plt.tight_layout()
plt.show()

→ Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz
11490434/11490434 0s 0us/step
Creating VAE model...
Training VAE...
Epoch 1/5
469/469 18s 27ms/step - kl_loss: 23.6851 - loss: 370.3940 - reconstruction_loss: 346.7090
Epoch 2/5
469/469 19s 25ms/step - kl_loss: 13.0540 - loss: 208.6858 - reconstruction_loss: 195.6317
Epoch 3/5
469/469 9s 19ms/step - kl_loss: 8.3095 - loss: 194.5787 - reconstruction_loss: 186.2693
Epoch 4/5
469/469 10s 21ms/step - kl_loss: 6.5814 - loss: 189.7816 - reconstruction_loss: 183.2001
Epoch 5/5
469/469 10s 22ms/step - kl_loss: 5.7807 - loss: 186.3691 - reconstruction_loss: 180.5884
Generating new digits...
1/1 0s 80ms/step

import tensorflow as tf
from tensorflow.keras.layers import Input, Dense, Lambda, Flatten, Reshape, Layer
from tensorflow.keras.models import Model
import numpy as np
import matplotlib.pyplot as plt

# Load and preprocess MNIST data
(x_train, _), (_, _) = tf.keras.datasets.mnist.load_data()
x_train = x_train.astype('float32') / 255.
x_train = x_train.reshape((len(x_train), 28, 28, 1))

# Network parameters
input_shape = (28, 28, 1)
batch_size = 128
latent_dim = 2
intermediate_dim = 512
epochs = 10

class Sampling(Layer):
    def call(self, inputs):
        z_mean, z_log_var = inputs
        batch = tf.shape(z_mean)[0]
        dim = tf.shape(z_mean)[1]
        epsilon = tf.random.normal(shape=(batch, dim))
        return z_mean + tf.exp(0.5 * z_log_var) * epsilon

# Encoder
encoder_inputs = Input(shape=input_shape)
x = Flatten()(encoder_inputs)
x = Dense(intermediate_dim, activation='relu')(x)
z_mean = Dense(latent_dim, name='z_mean')(x)
z_log_var = Dense(latent_dim, name='z_log_var')(x)
z = Sampling()([z_mean, z_log_var])
encoder = Model(encoder_inputs, [z_mean, z_log_var, z], name='encoder')

# Decoder
decoder_inputs = Input(shape=(latent_dim,))
x = Dense(intermediate_dim, activation='relu')(decoder_inputs)
x = Dense(784, activation='sigmoid')(x)
decoder_outputs = Reshape((28, 28, 1))(x)
decoder = Model(decoder_inputs, decoder_outputs, name='decoder')

class VAE(tf.keras.Model):
    def __init__(self, encoder, decoder, **kwargs):
        super().__init__(**kwargs)
        self.encoder = encoder
        self.decoder = decoder
        self.total_loss_tracker = tf.keras.metrics.Mean(name='total_loss')
        self.reconstruction_loss_tracker = tf.keras.metrics.Mean(name='reconstruction_loss')
        self.kl_loss_tracker = tf.keras.metrics.Mean(name='kl_loss')

    @property
    def metrics(self):
        return [
            self.total_loss_tracker,
            self.reconstruction_loss_tracker,
            self.kl_loss_tracker
        ]

    def train_step(self, data):
        ...

```

```

with tf.GradientTape() as tape:
    # Encoder output
    z_mean, z_log_var, z = self.encoder(data)
    # Reconstruction
    reconstruction = self.decoder(z)

    # Flatten input and reconstruction for binary crossentropy
    flat_data = tf.reshape(data, (-1, 784))
    flat_reconstruction = tf.reshape(reconstruction, (-1, 784))

    # Reconstruction loss (fixed axis handling)
    reconstruction_loss = tf.reduce_mean(
        tf.keras.losses.binary_crossentropy(
            flat_data,
            flat_reconstruction
        )
    ) * 784

    # KL divergence loss
    kl_loss = -0.5 * tf.reduce_mean(
        tf.reduce_sum(
            1 + z_log_var - tf.square(z_mean) - tf.exp(z_log_var),
            axis=1
        )
    )

    # Total loss
    total_loss = reconstruction_loss + kl_loss

# Compute gradients
grads = tape.gradient(total_loss, self.trainable_weights)
self.optimizer.apply_gradients(zip(grads, self.trainable_weights))

# Update metrics
self.total_loss_tracker.update_state(total_loss)
self.reconstruction_loss_tracker.update_state(reconstruction_loss)
self.kl_loss_tracker.update_state(kl_loss)

return {
    "loss": self.total_loss_tracker.result(),
    "reconstruction_loss": self.reconstruction_loss_tracker.result(),
    "kl_loss": self.kl_loss_tracker.result()
}

def call(self, inputs):
    z_mean, z_log_var, z = self.encoder(inputs)
    reconstruction = self.decoder(z)
    return reconstruction

# Create and compile VAE
print("Creating VAE model...")
vae = VAE(encoder, decoder)
vae.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=1e-4))

# Train VAE
print("Training VAE...")
vae.fit(x_train,
        epochs=epochs,
        batch_size=batch_size,
        shuffle=True)

def generate_digits(decoder, n=10):
    """Generate n new digits using the decoder"""
    z_sample = np.random.normal(size=(n, latent_dim))
    x_decoded = decoder.predict(z_sample)
    return x_decoded

# After training the VAE, you can use it like this:
print("\nGenerating new digits...")
new_digits = generate_digits(decoder, 10) # Generates 10 new digits

# To visualize the generated digits
plt.figure(figsize=(10, 2))
for i in range(10):
    plt.subplot(1, 10, i + 1)
    plt.imshow(new_digits[i].reshape(28, 28), cmap='gray')
    plt.axis('off')
plt.tight_layout()
plt.show()

```

Creating VAE model...
 Training VAE...
 Epoch 1/10
 469/469 12s 21ms/step - kl_loss: 22.7001 - loss: 373.2144 - reconstruction_loss: 350.5143
 Epoch 2/10
 469/469 10s 20ms/step - kl_loss: 12.1085 - loss: 212.6563 - reconstruction_loss: 200.5479
 Epoch 3/10
 469/469 10s 22ms/step - kl_loss: 8.1355 - loss: 196.1914 - reconstruction_loss: 188.0559
 Epoch 4/10
 469/469 10s 22ms/step - kl_loss: 6.6512 - loss: 190.8534 - reconstruction_loss: 184.2022
 Epoch 5/10
 469/469 10s 22ms/step - kl_loss: 6.0624 - loss: 186.8663 - reconstruction_loss: 180.8038
 Epoch 6/10
 469/469 9s 20ms/step - kl_loss: 5.6774 - loss: 183.2273 - reconstruction_loss: 177.5500
 Epoch 7/10
 469/469 11s 20ms/step - kl_loss: 5.3752 - loss: 180.2263 - reconstruction_loss: 174.8511
 Epoch 8/10
 469/469 10s 22ms/step - kl_loss: 5.1748 - loss: 177.3669 - reconstruction_loss: 172.1921
 Epoch 9/10
 469/469 10s 22ms/step - kl_loss: 5.0328 - loss: 174.8165 - reconstruction_loss: 169.7837
 Epoch 10/10
 469/469 10s 22ms/step - kl_loss: 4.9197 - loss: 172.7376 - reconstruction_loss: 167.8179



```

import tensorflow as tf
from tensorflow.keras.layers import Input, Dense, Lambda, Flatten, Reshape, Layer
from tensorflow.keras.models import Model
import numpy as np

```

7/5/25, 1:52 PM

```

import matplotlib.pyplot as plt

# Load and preprocess MNIST data
(x_train, _), (_, _) = tf.keras.datasets.mnist.load_data()
x_train = x_train.astype('float32') / 255.
x_train = x_train.reshape((len(x_train), 28, 28, 1))

# Network parameters
input_shape = (28, 28, 1)
batch_size = 128
latent_dim = 2
intermediate_dim = 512
epochs = 30

class Sampling(Layer):
    def call(self, inputs):
        z_mean, z_log_var = inputs
        batch = tf.shape(z_mean)[0]
        dim = tf.shape(z_mean)[1]
        epsilon = tf.random.normal(shape=(batch, dim))
        return z_mean + tf.exp(0.5 * z_log_var) * epsilon

# Encoder
encoder_inputs = Input(shape=input_shape)
x = Flatten()(encoder_inputs)
x = Dense(intermediate_dim, activation='relu')(x)
z_mean = Dense(latent_dim, name='z_mean')(x)
z_log_var = Dense(latent_dim, name='z_log_var')(x)
z = Sampling([z_mean, z_log_var])
encoder = Model(encoder_inputs, [z_mean, z_log_var, z], name='encoder')

# Decoder
decoder_inputs = Input(shape=(latent_dim,))
x = Dense(intermediate_dim, activation='relu')(decoder_inputs)
x = Dense(784, activation='sigmoid')(x)
decoder_outputs = Reshape((28, 28, 1))(x)
decoder = Model(decoder_inputs, decoder_outputs, name='decoder')

class VAE(tf.keras.Model):
    def __init__(self, encoder, decoder, **kwargs):
        super().__init__(**kwargs)
        self.encoder = encoder
        self.decoder = decoder
        self.total_loss_tracker = tf.keras.metrics.Mean(name='total_loss')
        self.reconstruction_loss_tracker = tf.keras.metrics.Mean(name='reconstruction_loss')
        self.kl_loss_tracker = tf.keras.metrics.Mean(name='kl_loss')

    @property
    def metrics(self):
        return [
            self.total_loss_tracker,
            self.reconstruction_loss_tracker,
            self.kl_loss_tracker
        ]

    def train_step(self, data):
        with tf.GradientTape() as tape:
            # Encoder output
            z_mean, z_log_var, z = self.encoder(data)
            # Reconstruction
            reconstruction = self.decoder(z)

            # Flatten input and reconstruction for binary crossentropy
            flat_data = tf.reshape(data, (-1, 784))
            flat_reconstruction = tf.reshape(reconstruction, (-1, 784))

            # Reconstruction loss (fixed axis handling)
            reconstruction_loss = tf.reduce_mean(
                tf.keras.losses.binary_crossentropy(
                    flat_data,
                    flat_reconstruction
                )
            ) * 784

            # KL divergence loss
            kl_loss = -0.5 * tf.reduce_mean(
                tf.reduce_sum(
                    1 + z_log_var - tf.square(z_mean) - tf.exp(z_log_var),
                    axis=1
                )
            )

            # Total loss
            total_loss = reconstruction_loss + kl_loss

        # Compute gradients
        grads = tape.gradient(total_loss, self.trainable_weights)
        self.optimizer.apply_gradients(zip(grads, self.trainable_weights))

        # Update metrics
        self.total_loss_tracker.update_state(total_loss)
        self.reconstruction_loss_tracker.update_state(reconstruction_loss)
        self.kl_loss_tracker.update_state(kl_loss)

        return {
            "loss": self.total_loss_tracker.result(),
            "reconstruction_loss": self.reconstruction_loss_tracker.result(),
            "kl_loss": self.kl_loss_tracker.result()
        }

    def call(self, inputs):
        z_mean, z_log_var, z = self.encoder(inputs)
        reconstruction = self.decoder(z)
        return reconstruction

# Create and compile VAE
print("Creating VAE model...")
vae = VAE(encoder, decoder)
vae.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=1e-4))

# Train VAE
print("Training VAE...")
vae.fit(x_train,
        epochs=epochs,
        batch_size=batch_size)

```

```

batch_size=batch_size,
shuffle=True)

def generate_digits(decoder, n=10):
    """Generate n new digits using the decoder"""
    z_sample = np.random.normal(size=(n, latent_dim))
    x_decoded = decoder.predict(z_sample)
    return x_decoded

# After training the VAE, you can use it like this:
print("\nGenerating new digits...")
new_digits = generate_digits(decoder, 10) # Generates 10 new digits

# To visualize the generated digits
plt.figure(figsize=(10, 2))
for i in range(10):
    plt.subplot(1, 10, i + 1)
    plt.imshow(new_digits[i].reshape(28, 28), cmap='gray')
    plt.axis('off')
plt.tight_layout()
plt.show()

Creating VAE model...
Training VAE...
Epoch 1/30
469/469 12s 22ms/step - kl_loss: 20.3889 - loss: 392.4920 - reconstruction_loss: 372.1030
Epoch 2/30
469/469 10s 22ms/step - kl_loss: 11.2360 - loss: 214.4868 - reconstruction_loss: 203.2508
Epoch 3/30
469/469 10s 22ms/step - kl_loss: 8.2049 - loss: 198.2868 - reconstruction_loss: 190.0818
Epoch 4/30
469/469 9s 19ms/step - kl_loss: 6.9466 - loss: 191.6655 - reconstruction_loss: 184.7189
Epoch 5/30
469/469 11s 20ms/step - kl_loss: 6.1922 - loss: 187.7495 - reconstruction_loss: 181.5573
Epoch 6/30
469/469 10s 22ms/step - kl_loss: 5.7655 - loss: 185.0690 - reconstruction_loss: 179.3036
Epoch 7/30
469/469 10s 21ms/step - kl_loss: 5.4850 - loss: 182.2723 - reconstruction_loss: 176.7873
Epoch 8/30
469/469 10s 22ms/step - kl_loss: 5.2473 - loss: 179.2172 - reconstruction_loss: 173.9700
Epoch 9/30
469/469 9s 19ms/step - kl_loss: 5.0517 - loss: 176.5014 - reconstruction_loss: 171.4497
Epoch 10/30
469/469 11s 21ms/step - kl_loss: 4.9018 - loss: 174.7234 - reconstruction_loss: 169.8217
Epoch 11/30
469/469 10s 22ms/step - kl_loss: 4.8123 - loss: 173.2358 - reconstruction_loss: 168.4235
Epoch 12/30
469/469 11s 23ms/step - kl_loss: 4.7895 - loss: 172.1038 - reconstruction_loss: 167.3143
Epoch 13/30
469/469 19s 21ms/step - kl_loss: 4.7976 - loss: 170.4304 - reconstruction_loss: 165.6328
Epoch 14/30
469/469 10s 20ms/step - kl_loss: 4.8200 - loss: 168.9783 - reconstruction_loss: 164.1583
Epoch 15/30
469/469 11s 22ms/step - kl_loss: 4.8575 - loss: 168.4328 - reconstruction_loss: 163.5752
Epoch 16/30
469/469 10s 22ms/step - kl_loss: 4.8760 - loss: 167.5619 - reconstruction_loss: 162.6859
Epoch 17/30
469/469 10s 22ms/step - kl_loss: 4.8580 - loss: 167.1993 - reconstruction_loss: 162.3413
Epoch 18/30
469/469 10s 21ms/step - kl_loss: 4.9139 - loss: 166.1452 - reconstruction_loss: 161.2314
Epoch 19/30
469/469 9s 20ms/step - kl_loss: 4.9181 - loss: 165.8288 - reconstruction_loss: 160.9108
Epoch 20/30
469/469 11s 22ms/step - kl_loss: 4.9585 - loss: 165.3715 - reconstruction_loss: 160.4131
Epoch 21/30
469/469 10s 22ms/step - kl_loss: 4.9537 - loss: 164.5709 - reconstruction_loss: 159.6171
Epoch 22/30
469/469 20s 22ms/step - kl_loss: 4.9938 - loss: 164.4708 - reconstruction_loss: 159.4770
Epoch 23/30
469/469 9s 20ms/step - kl_loss: 5.0328 - loss: 163.9709 - reconstruction_loss: 158.9381
Epoch 24/30
469/469 10s 20ms/step - kl_loss: 5.0645 - loss: 163.2563 - reconstruction_loss: 158.1918
Epoch 25/30
469/469 10s 22ms/step - kl_loss: 5.0860 - loss: 163.0060 - reconstruction_loss: 157.9200
Epoch 26/30
469/469 10s 22ms/step - kl_loss: 5.1121 - loss: 162.8665 - reconstruction_loss: 157.7544
Epoch 27/30
469/469 10s 22ms/step - kl_loss: 5.1063 - loss: 162.5496 - reconstruction_loss: 157.4433
Epoch 28/30
469/469 20s 20ms/step - kl_loss: 5.1355 - loss: 162.1430 - reconstruction_loss: 157.0075
Epoch 29/30
469/469 11s 22ms/step - kl_loss: 5.1699 - loss: 162.1090 - reconstruction_loss: 156.9391
Epoch 30/30
469/469 21s 22ms/step - kl_loss: 5.1833 - loss: 161.3323 - reconstruction_loss: 156.1491

Generating new digits...
1/1 0s 78ms/step

```



Task 3A: How the latent space changed with number of epochs (≈ 100 words)

As the number of training epochs increased, the 2D latent space became more structured and separated. With just 5 epochs, the latent space was loosely clustered, and digits were overlapping, indicating the model hadn't learned clear distinctions. At 15 epochs, the digits started forming tighter, more distinct clusters, showing that the model was better at grouping similar digits. By 30 epochs, the clusters became well-separated, suggesting that the encoder learned a meaningful, compressed representation of each digit. Overall, more training helped the model organize the latent space more efficiently and clearly.

Task 3B: How the quality of generated digits changed (≈ 100 words) The quality of generated digits significantly improved with more training epochs.

After 5 epochs, the generated digits were blurry and often unrecognizable. By 15 epochs, the digits became clearer and more shaped, though some still had irregular edges or overlapping strokes. At 30 epochs, the digits appeared sharp, well-formed, and closely resembled real MNIST digits. This progression shows that as the model saw more training examples, the decoder improved its ability to recreate realistic outputs from latent space points. Longer training directly contributed to more accurate and visually convincing digit generation.

